

<https://www.notion.so/Lecture-03-V8-Engine-Memory-Allocation-27a3a78e0e228017963adec0533f0d78>

Complete In-Depth Explanation of Memory Management in JavaScript

I'll explain this entire lecture comprehensively, covering every important concept without skipping anything.

Part 1: Foundation - How Programs Work

Understanding Basic Computer Architecture

When you have applications on your phone (PUBG, Instagram, WhatsApp), they're stored somewhere. Let me break this down:

1. Hard Disk/SSD (Permanent Storage)

- Your application code is stored here
- This is **permanent storage** - even when you switch off your device, data remains
- Example: Instagram's code lives here
- Speed: Relatively SLOW

2. RAM (Random Access Memory - Temporary Storage)

- When you want to RUN an application, the code is first loaded from Hard Disk → RAM
- RAM is **temporary storage** - when you close an app or switch off device, data is cleared
- Speed: Much FASTER than Hard Disk
- Why use RAM? Because transferring data from Hard Disk directly to CPU would be extremely slow

3. CPU (Central Processing Unit)

- The CPU executes the instructions from your code
- It reads instructions from RAM and performs operations

Flow Summary:

Hard Disk → RAM → CPU
(Code stored) (Code loaded) (Code executed)

Part 2: Understanding Memory Addressing

The Core Problem

Let's say you have:

```
a = 10  
b = 20
```

Both need to be stored in RAM. But here's the problem:

Scenario: Later you write:

```
b = 25 // Change b from 20 to 25
```

How does the computer know WHICH 20 to change? If there are multiple 20s in memory, how do we identify the correct one?

Solution: Memory Addressing

Think of RAM as having **blocks**, each of specific size (1 byte = 8 bits).

Example: 16-byte RAM

Let's divide this 16-byte RAM into blocks:

```
[Block 0][Block 1][Block 2][Block 3]...[Block 15]
```

Each block = 1 byte

We give each block a **unique address**:

- Block 0 → Address: 0
- Block 1 → Address: 1
- Block 2 → Address: 2
- And so on...

This is called **Byte Addressable** - every byte gets an address.

How Variables Map to Addresses

```
a = 10 // Let's say stored at address 2  
b = 20 // Let's say stored at address 8
```

Behind the scenes, your code is transformed:

```
// What you write:  
a = 10
```

```
// What actually happens:  
address_0010 = 10 // 'a' is replaced with its memory address
```

The Mapping:

- Variable **a** → Address 2 → Value 10
- Variable **b** → Address 8 → Value 20

When you update:

```
b = 25
```

The computer:

1. Looks up where **b** is stored (Address 8)
2. Goes to Address 8
3. Changes 20 → 25

No confusion! Because each variable has a unique address.

Part 3: Bit Requirements for Addresses

Address Size Calculation

For 16-byte RAM:

- Total bytes = 16
- Need 16 unique addresses (one per byte)
- To represent 16 different values, you need: **4 bits**
 - Because $2^4 = 16$

Address Representation:

Address 0: 0000

Address 1: 0001

Address 2: 0010

Address 3: 0011

...

Address 15: 1111

For 32-byte RAM:

- Need 32 addresses
- $2^5 = 32$, so need **5 bits**

Real-World Systems

32-bit Operating System:

- Address size = 32 bits = 4 bytes
- Can create 2^{32} different addresses
- Can handle up to **4 GB RAM**

64-bit Operating System:

- Address size = 64 bits = 8 bytes
 - Can create 2^{64} addresses
 - Can handle RAM > 4 GB (8 GB, 16 GB, 32 GB, etc.)
-

Part 4: Stack vs Heap - The Genesis

Now comes the crucial part: **Why do we need TWO types of memory (Stack and Heap)?**

Initial Approach: Sequential Storage

Simple Strategy:

RAM: [_____]

Store data sequentially:

a = 10 → [10][][][]...

b = 20 → [10][20][][]...

c = 30 → [10][20][30][]...

Rule: Next data goes right after the previous one.

This is essentially **STACK** - data stored one after another (Last In, First Out).

The Problem Arises

Now, what if:

```
b = "Sun" // Change b from 20 to a string
```

Issue:

- The number **20** takes 1 byte
- The string "**Sun**" needs 3 bytes (S + U + N, each character = 1 byte)
- The original space (1 byte) is NOT enough!

What are the options?

Option 1: Shift everything

Before: [10][20][30][][]...

After: [10][S][U][N][30]...

Problems with Option 1:

1. **Slow:** Need to move all subsequent data
2. **Address Changes:** The address of **30** changed from position 3 to position 5
3. If **c** is used 50 times in code, all 50 references need address updates!

This is **extremely inefficient**.

Alternative Approach: Random Allocation (Heap)

Instead of sequential storage, store wherever space is available:

RAM: [_____]

a = 10 → Store at position 5

b = 20 → Store at position 12

c = 30 → Store at position 3

Benefits:

- Flexible space allocation
- Can allocate any amount of space needed

Same Problem Persists: When **b** changes from **20** to "Sun":

1. Need to find NEW space for 3 bytes
 2. Address of **b** changes
 3. All code references need updating
-

Part 5: The Brilliant Solution - Pointer System

The Insight

The lecturer had a genius realization:

Problem: When data changes location, addresses change, causing widespread code updates.

Solution: Use an **extra layer of indirection!**

How It Works

Step 1: Store actual data in HEAP

"Sun" → Stored at Heap address 0x1010

Step 2: Store the HEAP address in STACK

Variable b → Stack address 0x0004

Value at 0x0004 = 0x1010 (points to heap)

The Magic

When **b** changes:

b = "Mohit" // 5 characters now

Process:

1. Find new space in Heap for "Mohit" → say address 0x2020
2. **Only update the Stack:** Change 0x0004 from `0x1010` to `0x2020`

Result:

- Stack address of `b` (0x0004) **never changes**
- Code doesn't need updates
- Only the heap address pointer is modified

Two-Step Access

To access `b`:

1. Go to Stack (`b`'s address) → Get heap address
2. Go to that Heap address → Get actual data

Summary of Stack vs Heap

STACK:

- Stores addresses/pointers
- Fixed, small size (KBs to few MBs)
- Fast access
- Sequential allocation

HEAP:

- Stores actual data
- Large size (MBs to GBs)
- Slower access (need to find free space)
- Random allocation

Part 6: Fixed-Size vs Dynamic-Size Data

Fixed-Size Data → Can Use Stack Directly?

For numbers like `10`, `20`, `30`:

- Size is fixed (always takes same space)
- If `a = 10` changes to `a = 50`, both take same space
- Could theoretically store directly in stack

BUT WAIT! There's a catch with JavaScript...

Part 7: JavaScript's Immutability Rule

Primitive Data Types Are IMMUTABLE

In JavaScript, **primitive values cannot be changed**.

```
let a = 10
```

When stored:

Memory location 0xABCD: [10]

Now:

```
a = 50
```

JavaScript's Rule: Cannot modify the original **10**. Instead:

1. Create NEW memory location
2. Store **50** there (say address 0xDEF)
3. Update **a** to point to new address

The Problem This Creates

Even though numbers are fixed-size, **their memory location changes** due to immutability!

Consequence:

```
let a = 10 // Address: 0xABCD
```

// Later:

```
a = 50 // New address: 0xDEF
```

If **a** is used 50 times in code, all 50 references need address update from **0xABCD** to **0xDEF**.

This is slow!

Part 8: The Final Solution - Everything in Heap

JavaScript's Approach

ALL DATA goes to Heap (primitives AND non-primitives), but addresses stored in Stack.

```
let a = 10  
let b = 20
```

Memory Layout:

HEAP:

Address 0x1000: [10]
Address 0x1008: [20]

STACK:

Variable a → 0x1000 (points to heap)
Variable b → 0x1008 (points to heap)

When Value Changes

```
a = 50
```

Process:

1. Create new location in Heap → 0x2000
2. Store 50 there
3. Update Stack: a now points to 0x2000 (instead of 0x1000)

Key Point: Stack's address for variable a **never changes**. Only the heap address it points to changes.

Part 9: Optimization 1 - Pre-allocated Values

The Special Cases

For true, false, null, undefined:

- Only a few possible values
- Used very frequently

Optimization: When program starts, **pre-allocate** these in Heap:

```
Heap Address 0x0001: true  
Heap Address 0x0002: false  
Heap Address 0x0003: null  
Heap Address 0x0004: undefined
```

Benefit: All boolean variables can point to **same true or false**:

```
let x = true // Points to 0x0001  
let y = true // Also points to 0x0001  
let z = true // Also points to 0x0001
```

No need to create multiple **true** values!

Part 10: The Loop Performance Problem

The Inefficiency

```
for(let i = 0; i < 100; i++) {  
    console.log(i)  
}
```

What happens:

- $i = 0 \rightarrow$ Create in Heap, allocate memory
- $i = 1 \rightarrow$ Create NEW location in Heap, allocate memory
- $i = 2 \rightarrow$ Create NEW location in Heap, allocate memory
- ...continues 100 times

Problem:

- Finding free space in Heap is **SLOW**
- Allocating memory 100 times makes loop **very slow**

Attempted Solution: Pre-allocated Array

Create an array with numbers 0 to 100,000 when program starts:

```
const numbers = [0, 1, 2, 3, ..., 100000]
```

Memory:

Base Address: 1000

numbers[0] at 1000

numbers[1] at 1008 (1000 + 8 bytes)

numbers[2] at 1016 (1000 + 8*2)

...

Formula to access:

Address = Base Address + (Index × Size of Data)

Benefit: No need to search for free space!

Problem: Uses $100,000 \times 8$ bytes = 800 KB just for this array!

Part 11: The Genius Solution - SMI (Small Integers)

The Breakthrough

Address size in 32-bit system = 32 bits (4 bytes)

Key Insight: Use the address itself to store small numbers!

How It Works

An address looks like:

[31 bits][1 bit]

Strategy:

- **Last bit = 0:** The address contains a number directly
- **Last bit = 1:** The address points to heap data

Example

let i = 5

Instead of:

Stack: i → 0x1000

Heap: 0x1000 → [5]

Do:

Stack: i → [0000...0101][0]

↑ ↑
number 5 flag=0 (is number)

The first 31 bits **store the number directly!**

Reading the Value

To get the value:

```
// If address = 0x00000028 (binary: ...0101000)  
// Last bit = 0 → It's a number  
// First 31 bits = 5
```

```
actual_value = address >> 1 // Shift right by 1 bit
```

Range of SMI

With 31 bits:

- Total combinations = 2^{31}
- Half for positive: 0 to $2^{30} - 1$
- Half for negative: -2^{30} to -1

First bit determines sign:

- First bit = 0 → Positive
- First bit = 1 → Negative

Benefits

1. **No heap allocation needed** for small numbers
2. **No memory search** required
3. **Instant value access** - just read the address
4. **Loop performance:** Dramatically faster!

```
for(let i = 0; i < 100; i++)
```

Now just increments address value directly - **no heap operations!**

Part 12: Complete Memory Model

Final Architecture

STACK:

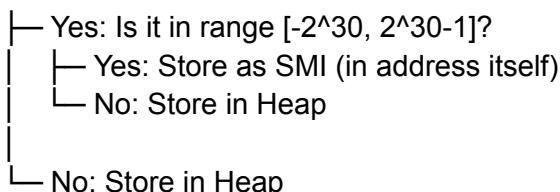
- Stores addresses only
- Address size fixed (32-bit or 64-bit)

HEAP:

- Stores actual data
- All dynamic/large data
- Numbers outside SMI range
- Strings, objects, arrays

Decision Tree

Is it a number?



Why Everything Appears in Heap

Technically:

- Small integers (SMI): Stored as addresses (not actually in heap)
- Large numbers: Stored in Heap
- Strings: Stored in Heap

- Objects: Stored in Heap
- Arrays: Stored in Heap

For simplicity, we say: "Everything goes to heap" because most data does, and SMI is an optimization detail.

Part 13: Garbage Collection

The Problem

```
let a = 10 // Heap: 0x1000 → [10]
a = 50    // Heap: 0x2000 → [50]
```

Now **[10]** at 0x1000 is **orphaned** - no variable points to it.

Solution: Garbage Collector

Process:

1. **Mark Phase:** Start from Stack, mark all reachable heap data
2. **Sweep Phase:** Delete all unmarked heap data

Example:

Stack:

```
a → 0x2000
b → 0x3000
```

Heap:

```
0x1000: [10] ← Orphaned
0x2000: [50] ← Marked (a points here)
0x3000: [20] ← Marked (b points here)
```

Garbage collector removes data at 0x1000.

Why Important:

- Frees up heap space
 - Prevents memory leaks
 - Prevents program crashes from running out of memory
-

Part 14: 64-bit Systems

In 64-bit systems:

- Address size = 64 bits = 8 bytes
 - More bits available for SMI
 - Can represent even larger range of small integers
-

Summary of Key Concepts

1. **Memory Hierarchy:** Hard Disk → RAM → CPU
 2. **Addressing:** Each byte gets unique address
 3. **Stack vs Heap:** Stack for addresses, Heap for data
 4. **Pointer System:** Extra indirection layer prevents widespread address updates
 5. **Immutability:** JavaScript primitives can't be modified in place
 6. **SMI Optimization:** Store small integers in address itself
 7. **Garbage Collection:** Automatic cleanup of unreferenced data
-

This system demonstrates brilliant engineering - solving performance problems through clever bit manipulation and understanding of computer architecture!

1. The Big Picture: Where Does Your Code and Data Live? (Permanent vs. Temporary Storage)

First Principle: Computers are machines that execute instructions on data. Instructions (code) and data must be *stored somewhere* and *accessed quickly*. Storage speed vs. capacity is the core trade-off.

- **Permanent Storage (Disk/SSD/HDD):** Your apps (PUBG, Instagram) are files on disk.
Why? Disks are cheap, high-capacity (TB+), but *slow* (ms latency). Code here is "dormant"—like a book on a shelf.
 - Example: Instagram.apk (or .exe) is ~100MB on your SSD.
 - Persistence: Survives power-off. Switch off your phone? Apps don't vanish.
- **Temporary Storage (RAM - Random Access Memory):** To run code, load it into RAM.
Why RAM? *Ultra-fast* access (ns latency), but volatile (loses data on power-off) and expensive (~\$50/GB).
 - Process: OS (Android/iOS/Windows) copies code from disk → RAM. CPU fetches instructions from RAM, executes them.

- Why not direct disk-to-CPU? Disk is 10^6 x slower than RAM. Bottleneck: CPU idles waiting for data.
- Analogy: Disk = library (vast, slow walk). RAM = desk (small, instant grab). CPU = you reading at the desk.

Aspect	Permanent Storage (Disk/SSD)	Temporary Storage (RAM)
Speed	Slow (10-100ms seek time)	Fast (10-100ns access)
Capacity	High (GB-TB, cheap)	Low (GB, expensive)
Persistence	Yes (survives reboot)	No (cleared on shutdown)
Use Case	Store code/data at rest	Execute code/data in flight
Size	1TB SSD for 1000s of apps	8GB RAM for 5-10 active apps
Example		

Derivation: RAM size limits multitasking. 16GB RAM? Run 2-3 heavy apps (PUBG + Instagram) before swapping to disk (thrashing → lag). Close PUBG? OS deallocates its RAM chunk.

Pitfall: Forgetting this leads to "memory leaks" later—unused data clogs RAM, forcing swaps.

2. Inside RAM: Byte-Addressable Memory and Unique Identification

First Principle: RAM is a sea of bits. To store/retrieve data, we need *addresses*—like house numbers. Without them, "where is my 20?" becomes a search nightmare.

- **Byte-Addressable Memory:** RAM divided into 8-bit chunks (bytes). Each byte gets a unique address (0 to N-1).
 - Why bytes? 1 byte = 256 values (2^8). Fits chars (ASCII), small nums. Larger data (e.g., int) spans multiple bytes.

Example: 16-byte RAM (toy model). Addresses: 0x00 to 0x0F (hex for brevity).

text

RAM Layout (16 bytes):

Addr: 00 01 02 03 04 05 06 07 08 09 0A 0B 0C 0D 0E 0F

- Byte: [] [] [10][] [20][] [] [] [] [] [] [] [] []
 - Store a = 10: Assume 10 fits 1 byte (binary: 00001010). Put at addr 02.
 - Store b = 20: Binary 00010100 at addr 04.
- **The Problem: Uniqueness:** Update b = 25. Scan RAM for "20"? Multiple 20s? Data corruption!

- Solution: **Mapping Table** (implicit in code). Variables → Addresses.
 - a maps to 0x02 (holds 10).
 - b maps to 0x04 (holds 20).

Behind scenes: Compiler/OS replaces a with its address (0x02). Code becomes: "Load from 0x02, store 10."

text

Source Code: let a = 10; b = 20; b = 25;

- Compiled View: [addr:0x02] = 10; [addr:0x04] = 20; [addr:0x04] = 25;
- **Address Size Math:** For M bytes, need $\log_2(M)$ bits to address them.
 - 16 bytes: $2^4 = 16 \rightarrow 4$ bits (0000 to 1111).
 - 32 bytes: $2^5 = 32 \rightarrow 5$ bits.
 - Real RAM (4GB = 2^{32} bytes): 32 bits (4 bytes/address).
 - 64-bit system: 64 bits (8 bytes/address) → Addresses 16 exabytes (huge).

RAM Size	Addresses Needed	Bits per Address	Binary Example
16 bytes	16	4	0000 to 1111
1KB (1024B)	1024	10	0000000000+
4GB (2^{32} B)	2^{32}	32	0x00000000+

Derivation: Addresses stored in binary. For 32-bit: 4 bytes/address. Total addressable: 2^{32} bytes = 4GB. Exceed? OS uses virtual memory (pages to disk).

Pitfall: Fixed address size limits RAM. 32-bit OS? Max ~4GB user RAM (even if hardware has more).

3. Stack vs. Heap: Sequential vs. Flexible Allocation

First Principle: Allocation strategies trade speed for flexibility. Fixed-size data? Predictable. Dynamic? Free-form.

- **The Split:** RAM divided: Stack (small, fast, sequential) + Heap (large, flexible, scattered).
 - Stack: LIFO (Last In, First Out). Grows/shrinks predictably (e.g., function calls).
 - Size: KB-MB (e.g., 1MB default).

Allocation: Sequential. Next item = last + fixed size.

text

Stack Growth (downward, convention):

```

High Addr
| ... |
| c=30 (addr 0x10, 1 byte) |
| b=20 (addr 0x09, 1 byte) |
| a=10 (addr 0x02, 1 byte) | ← Top (growing down)

```

- Low Addr
- Pros: Fast (no search). Fixed-size data (ints: 8 bytes) fits perfectly.
- Cons: Variable-size (strings) breaks it—overwrite neighbors or shift everything ($O(n)$ time, address updates everywhere).
- Heap: Dynamic. Allocate anywhere free.
 - Size: MB-GB (rest of RAM).

Allocation: Search free blocks (e.g., buddy system). References (addresses) stored in stack.
text

```

Heap (scattered):
Addr: 0x1000 [Rohit (5 bytes: R-O-H-I-T)] ← Dynamic string
0x2000 [free]
0x3000 [Array: [1,2,3] (24 bytes)]

```

- Stack Ref: d → 0x1000 (8-byte address)

Feature	Stack	Heap
Allocation	Sequential (next slot)	Scattered (find free block)
Size	Fixed per item (predictable)	Variable (dynamic)
Speed	Ultra-fast (no search)	Slower (search + fragmentation)
Data Types	Primitives (ints, bools)	Objects, arrays, strings
Access	Direct value	Via reference (2-step: stack → heap)
Deallocatio n	Auto (scope end)	Manual/GC (later)

- **Why Split?** From scratch: Sequential for small/fixed (no waste). Flexible for big/variable (no shifts).
 - Fixed-size update: a=50 → Overwrite slot ($O(1)$).
 - Variable: b="Rohit" (5 bytes) → Can't fit 1-byte slot. Shift? Update 50+ code refs (slow). Solution: Heap + ref in stack.
 - Math: String "Rohit" = 5 chars × 1 byte/char (UTF-8 basic) = 5 bytes + null terminator.

Derivation: Fragmentation in heap: Free blocks scatter → "holes." Allocators merge/coalesce.

Pitfall: Stack overflow (deep recursion). Heap fragmentation → OOM (Out of Memory).

4. JavaScript Specifics: Everything in Heap, Immutability, and References

First Principle: JS is high-level—hides low-level details. But V8 (JS engine) must map to machine code efficiently.

- **Immutability of Primitives:** Primitives (number, string, boolean, null, undefined, symbol) can't mutate. Why? Thread-safety, predictability.
 - let a = 10; a = 20; → Not mutate 10. Create new 20, re-point a.
 - Generic: Even fixed-size primitives go to heap (avoids address churn).

All in Heap: Variables (in stack-like "activation records") hold *references* (addresses, 8 bytes on 64-bit).

text

JS Code: let a = 10; let b = "Hi";

V8 Memory:

Stack: a → 0x7FFF... (ref to heap)

b → 0x7FFE... (ref to heap)

Heap: 0x7FFF... [10 (8 bytes, Number object)]

- 0x7FFE... ["Hi" (3 bytes + overhead)]
- **Access:** 2-step indirection. console.log(a) → Stack ref → Heap value.
 - Update: a=20 → New heap slot for 20, update stack ref. Old 10 orphaned → GC later.
- **Pre-Allocation for Constants:** true/false/null/undefined created at startup, single instances.
 - Benefit: Reuse refs (e.g., all true point to one heap slot). Saves space (no 100x true allocations).

JS Type	Mutable ?	Heap Allocation	Stack Holds	Example Update Cost
Number (primitive)	No	Yes (8 bytes)	Ref (8 bytes)	New alloc + ref swap
String	No	Yes (var len)	Ref	New alloc (copy chars)
Object	Yes	Yes (props dynamic)	Ref	Mutate in-place if space; else realloc
Array	Yes	Yes (elements)	Ref	Resize/realloc if overflow

Derivation: Ref size = pointer size (64-bit: 8 bytes). Total for a=10: 8 (ref) + 8 (value) = 16 bytes.

Pitfall: Immutability illusion—strings concat creates new strings (inefficient loops).

5. Garbage Collection: Cleaning the Heap

First Principle: Heap grows forever without cleanup → OOM crash.

- **Mark-and-Sweep GC:** V8's algorithm.
 - **Mark:** From stack roots (variables), follow refs. Mark reachable (e.g., a → 10 marked "live").
 - **Sweep:** Unmarked → Free (e.g., old 10 after a=20).
 - Frequency: Generational (young/old gen). Young: Frequent, small. Old: Rare, large.

text
Before GC:
Stack: a → Heap[20] (marked)
Heap: [10 (unmarked)] [20 (marked)] ["OldStr" (unmarked)]

After: Free [10] and "OldStr". Compact to reduce fragmentation.

- **Why Needed?** Recursion without base case: Stack fills (crash), but heap leaks if cycles.

Derivation: Time: O(heap size). Space: Compaction halves pauses.

Pitfall: GC pauses → Jank in UI. V8 mitigates with incremental GC.

6. Ultra-Optimization: Tagged Pointers for Small Integers (SMI)

First Principle: Most loop vars? Small ints (-2³⁰ to 2³⁰). Heap alloc per tick? Wasteful (search + alloc).

- **Problem in Loops:** for(let i=0; i<100; i++) → 100 heap allocs (slow).
 - Each: Find free slot (O(log N)), alloc 8 bytes.
- **Solution 1: Pre-Alloc Array (Naive).**
 - Startup: Create array[0...1e6] in heap (base addr 0x1000).
 - Access: i ref = base + i * 8.
 - Read: Heap[base + i*8].
 - But: Wastes 8MB upfront. Still indirection.
- **Solution 2: Virtual Array (No Real Alloc).**
 - Treat refs as math: Value = (ref - base) / 8.

- No heap visit for read—pure calc.
- **Solution 3: Tagged Pointers (V8 Magic).**
 - 64-bit pointer: 64 bits. Use LSB (bit 0) as tag.
 - LSB=0: SMI. Value = pointer $\gg 1$ (shift right, divide by 2).
 - LSB=1: Heap ref. Dereference for real data.
 - Why LSB? Alignments ensure pointers even (LSB=0 free).
 - Math:
 - Store 42: Binary ...0101010 | 0 (tag). Value = (...0101010 | 0) $\gg 1 = 42$.
 - Read: If LSB=0, $\gg 1 \rightarrow$ int. Else, load from addr.
 - Range: 31 bits signed (MSB sign bit): -2^{30} to $2^{30}-1$ (~1B).
 - Positive: 0xxx... \rightarrow Value as-is.
 - Negative: 1xxx... \rightarrow Two's complement.

Pointer Bits (64)	Tag (LSB)	Interpretation	Example (Hex)	Decoded Value
...0000 101010	0	SMI (small int)	0x...2A	42
...0001 000000	1	Heap ref (string/obj)	0x...1	Load[addr]
...1000 000000	0	SMI (negative)	0x...80	-64

- **Loop Win:** $i++ \rightarrow$ Increment pointer (add 2, keep LSB=0). No alloc/search.
 - $i=42 \rightarrow$ Pointer = $42 \ll 1 = 84$ (binary ...1010100).
 - Read: LSB=0 $\rightarrow 84 \gg 1 = 42$.

Derivation: 2^{63} total pointers. Half tagged (even). Signed: 2^{31} values.

Pitfall: Large ints ($>2^{30}$) \rightarrow Heap (slow). V8 detects, converts.

Conclusion: The Engineering Mindset

This isn't "stack=primitives, heap=objects" trivia—it's *why*: Efficiency from immutability + tagging. JS hides it, but understanding prevents leaks/slowness. Homework: Read V8 source on tagged pointers (or 64-bit variants—uses more bits). Replay: Build a toy allocator in Python.

Minds blown? That's the point. Next: Apply to leaks in your code. Jai Hind!